

# 9. Transaction Management

Page No.:

Date: .....

## Transaction concept

- i) Transaction is a set of logically related operations
- ii) A transaction is a unit of program execution that accesses & possibly updated various data items.
- iii) A transaction is delimited by statements of the form begin transaction & end transaction, the transaction consists of all operations executed between begin transaction & end transaction.

Ex) E.g. transaction to transfer \$50 from account

A to account B

read (A)

A = A - 50

write (A)

read (B)

B = B + 50

write (B)

## ACID Properties in DBMS

A = Atomicity  $\rightarrow$  The entire transaction takes place at once or doesn't happen at all.

C = consistency  $\rightarrow$  The database must be consistent before and after the transaction.

I = Isolation  $\rightarrow$  Multiple transactions occur independently without interference.

D = Durability  $\rightarrow$  The changes of successfully transaction occur even if the system failure occurs.

### ② Atomicity

i) Either the entire transaction takes place at once or doesn't happen at all. there is no midway i.e. transactions do not occur partially. Each

transaction is considered as one unit & either runs to completion or is not executed at all. It involves following two operations.

- ① Abort → If transaction aborts, changes made to db are not visible.
- ② commit → If transaction commits changes made are visible. Atomicity is also known as the "All or nothing rule."

e.g →

- consider two transaction  $T_1$  &  $T_2$ , Transfer of 100 from account  $X$  to account  $Y$ .

$T_1$	$T_2$
Before	
$X: 500$	$Y: 200$
Read( $X$ )	Read( $Y$ )
$X: X - 100$	$Y = Y + 100$
Write( $X$ )	Write( $Y$ )
$X: 400$	$Y = 300$

- If the transaction fails after completion of  $T_1$  but before completion of  $T_2$  then amount has been deducted in entirety in order to ensure correctness of DB state.

## 2) consistency

- i) This means that integrity constraints must be maintained so that DB is consistent before & after the transaction.
- ii) Referring to the above example, the total amount before & after the transaction must be maintained.

$$\text{Total before } T \text{ occurs} = 500 + 200 = 700$$

$$\text{Total after } T \text{ occurs} = 400 + 300 = 700$$

Therefore database is consistent. Inconsistency occurs in case  $T_1$  completes but  $T_2$  fails as a result  $T$  is incomplete.

### 3) Isolation

- i) This property ensures that multiple transactions can occur concurrently without leading to the inconsistency of DB state.
- ii) Transaction occur independently without interference.
- iii) changes occurring in a particular transaction will not be visible to any other transaction until that particular change in that trans. is written to memory or has been committed.
- iv) This property ensures that execution of trans. concurrently will result in a state i.e. equivalent to state achieved these were serially in some order, let  $x = 500$ ,  $y = 500$

e.g.  $\rightarrow T_1 \& T_2$

$T_1$	$T_2$
Read(x)	Read(x)
$x' = x * 100$	Read(y)
Write(x)	$z = x + y$
Read(y)	Write(z)
$y' = y - 60$	
Write	

- suppose  $T_1$  has been executed till Read(y) & then  $T_2$  start, as a result, interleaving of operations takes place due to which  $T_2$  reads correct value of x but incorrect value of y & sum computed by

$$T_2 = x + y = 50,000 + 600 = 50,500$$

is thus not consistent with the sum at end of trans.

$$T_1 = (x + y) = 50,000 + 4500 = 50,450$$

so result is inconsistency due to loss of 60 units.  
hence transaction must take

#### 4) Durability

- i) This property ensures that once the transaction has completed execution, the updates & modifications to the DB are stored in & written to disk & they persist even if a system failure occurs.
- ii) These updates now become permanent & are stored in non-volatile memory.

So the ACID properties in totality provide a mechanism to ensure correctness & consistency of DB in a way such that each trans. is a group of operations that acts as a single unit, produces consistent results, acts in DB isolation from other operations & updates that it makes are durably stored.

#### Transaction state

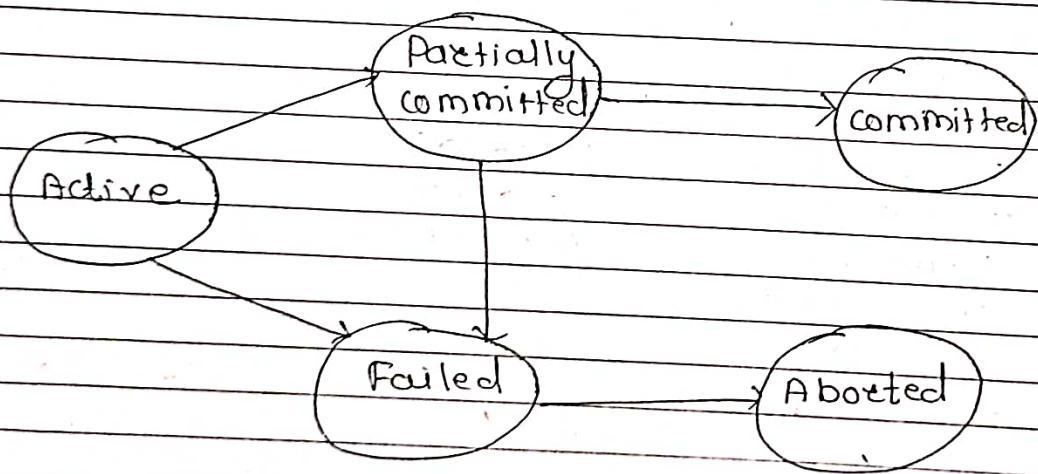
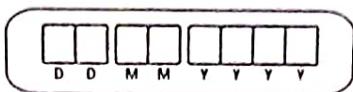


fig → state diagram of a transaction

#### 1) Active

- i) It is a initial state.
- ii) The transaction stays in this state while it is executing.
- iii) If transaction is in execution then it is said to be in active state.
- iv) It doesn't matter which step is in execution until

(6)



unless the transaction is executing, it remains in active state.

### Failed Partially committed

- i) In above fig. transaction goes into partially committed state from the active state when there are read & write operations present in transaction
- ii) A transaction contains no. of read & write opn once the whole trans. is successfully executed, the transaction goes into partially committed state where we have all the read & write operations performed on the main memory instead of the actual DB.
- iii) This state is used because transaction can fail during execution so if we are making changes in the actual DB instead of local memory, DB may be left in an inconsistent state in case of any failure.
- iv) This state helps us to rollback the changes made to the DB in case of failure during execution.

### 3) committed

- i) If transaction completes the execution successfully then all the changes made in the local memory during partially committed state, are permanently stored in DB
- ii) In above dia. transaction goes from partially committed state when everything is successful.

### 4) Aborted state

- i) If transaction fails during execution then the transaction goes into a failed state, the changes made into the local memory are rolled back to the previous consistent state & transaction goes into aborted state from the failure state.

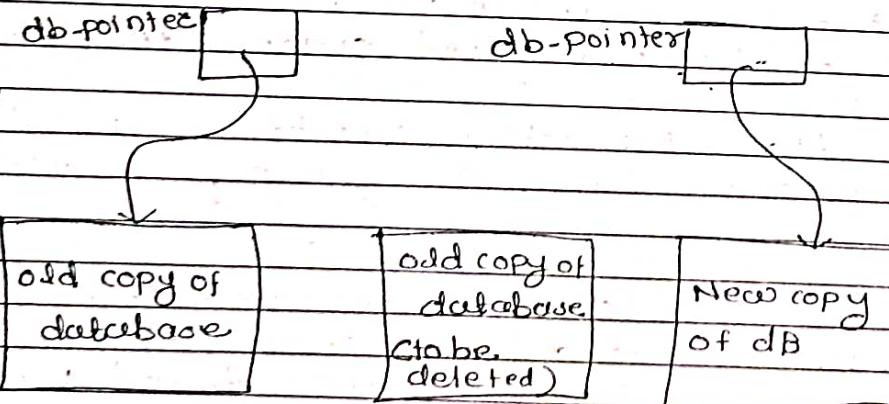


## Implementation of atomicity & Durability

- i) The recovery management component of DB system implements the support of atomicity & durability.

### Shadow Database scheme

- i) Assume that only one transaction is active at a time.
- ii) A pointer called db-pointer always points to the current consistent copy of the db.
- iii) All updates are made on a shadow copy of the DB & db-pointer is made to point to the updated shadow copy only after the transaction reaches partial commit & all updated pages have been flushed to disk.
- iv) In case transaction fails, old consistent copy pointed to by db-pointer can be used & the shadow copy can be deleted.



a) Before update

b) After update



## Transaction Isolation.

i) The transaction isolation is primarily used as means to provide accurate & reliable access to data within concurrent transactions.

ii) e.g. → Two different transactions might be accessing the same data simultaneously. Therefore, if the change made on the data by one transaction is not passed to other transaction, it can affect DB operations.

There are 4 different types of transaction isolation level.

1) Serializable → Implements read & write locks until the transaction is finished also implements range locks.

2) Repeatable reads

Implements read & write locks until the transaction is completed, doesn't manage range lock.

3) Read committed

Implements write lock until the transaction completed but releases read locks when SELECT operation is performed.

4) Read uncommitted

One transaction can see the uncommitted changes made by other transaction.

D	D	M	M	T	T	T
---	---	---	---	---	---	---

## Schedulability

### Types of schedules in DBMS

- i) Schedule is a process of defining the transaction & executing them one by one.
- ii) Schedule is a sequence of function instructions that specify the chronological order in which instructions of concurrent transactions are executed.
- iii) A schedule for a set of transaction must consist of all instructions of those transactions.
- iv) It must preserve the order in which the instructions appear in each individual transaction.
- v) A transaction that successfully completes its execution will have a commit instruction as the last statement.

### Schedule (1)

- Let T<sub>1</sub> transfer \$50 from A to B & T<sub>2</sub> transfer 20% of the balance from A to B
- A serial schedule in which T<sub>1</sub> is followed by T<sub>2</sub>:

T<sub>1</sub>                    T<sub>2</sub>

read(A)

A := A - 50

write(A)

read(B)

B := B + 50

write(B)

read(A)

temp := A \* 0.1

A := A - temp

write(A)

read(B)

B := B + temp

write(B)

Schedule ②

- A serial schedule where  $T_2$  is followed by  $T_1$

$T_1$	$T_2$
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
	read(B)
	B := B + temp
	write(B)
read(A)	
A := A - 50	
write(A)	
read(B)	
B := B + 50	
write(B)	

Schedule ③ (Non serial)

Let  $T_1$  &  $T_2$  be the transactions defined previously.

The following schedule is not a serial schedule.

but it is equivalent to schedule 1

$T_1$	$; T_2$
read(A)	
A := A - 50	
write(A)	
	read(A)
	temp := A * 0.1
	A := A - temp
	write(A)
read(B)	
B := B + 50	
write(B)	
	read(B)
	B := B + temp
	write(B)

In schedule 1, 2, & 3 the A+B preserved

### Schedule (4)

The following concurrent schedule does not preserve the value of  $(A+B)$

T <sub>1</sub>	T <sub>2</sub>
<code>read(A)</code>	<code>read(A)</code>
$A = A - 50$	$\text{temp} = A * 0.1$
	$A = A - \text{temp}$
<code>write(A)</code>	<code>write(B)</code>
<code>write(CB)</code>	
<code>write(CA)</code>	
<code>read(CB)</code>	
$B = B + 60$	
<code>write(CB)</code>	
	$B = B + \text{temp}$
	<code>write(B)</code>

### Serial scheduler

- i) schedules in which the transactions are executed i.e a serial schedule is one in which no transaction starts until a running transaction has ended are called serial schedules i.e. In serial schedule a transaction is executed completely before starting the execution of another transaction.
- ii) A transaction does not start execution until the currently running transaction finished execution.

### Non serial scheduler

- i) This is a type of scheduling where the operations of multiple transactions are interleaved.
- ii) The transactions are executed in a non serial manner, keeping the end result is correct & same as serial schedule.

In serial schedule where one trans. proceed without waiting for previous the non serial

- iii) In serial schedule where one transaction proceeds without waiting for the previous trans. to complete.

### Schedules

Serial Schedules

Non serial

Serializable schedule

Non-serializable schedule

conflict

view

recoverable

Non-recoverable

cascading

cascadeless

strict

### Serializability

- i) Serial schedule does not need the serializability because it follows transaction only when the previous transaction is complete.
- ii) The non serial schedule is said to be in serializable schedule only when it is equivalent to the serial schedules, for an no. of transactions.

These are two types

- 1) conflict
- 2) view

1) conflict serializability

i) If a schedule  $s$  can be transformed into a schedule  $s'$  by a series of swaps of non-conflicting instructions, we say that  $s$  &  $s'$  are conflict equivalent.

ii) We say that schedule  $s$  is conflict serializable if it is conflict equivalent to a serial schedule

iii) schedule ④ can be transformed into schedule ② a serial schedule where  $T_2$  follows  $T_1$  by series of swaps of non conflicting instructions, therefore schedule ④ is conflict serializable  
e.g →

$T_1$	$T_2$	$T_1$
read(A)		read(A)
write(A)		write(A)
	read(A)	Read(B)
	write(A)	write(B)
read(B)		
write(B)		
	read(B)	read(A)
	write(B)	write(A)

Schedule 1    Schedule 2

iv) E.g of a schedule that is not conflict serializable

$T_3$	$T_4$
read(g)	
write(g)	write(g)

v) we are enable to swap instruction in above schedule to obtain either the serial schedule  $\langle T_3, T_4 \rangle$  or serial schedule  $\langle T_4, T_3 \rangle$

## view serializability

- Let  $s$  &  $s'$  be two schedules with the same set of transactions.
- $s$  &  $s'$  are view equivalent if the following 3 conditions are met.
  - For each data item  $g$  if trans.  $T_1$  reads the initial value of  $g$  in schedule  $s$ , then trans.  $T_1$  must in schedule  $s'$  also read the initial value of  $g$ .
  - For each data item  $g$  if trans.  $T_1$  executes  $\text{read}(g)$  in schedule  $s$ , & that value was produced by trans.  $T_2$ , then trans.  $T_1$  must in schedule  $s'$  also read the value of  $g$  that was produced by trans.  $T_2$ .
  - For each data item  $g$ , the trans. that performs the final  $\text{write}(g)$  operation in schedule  $s$  must perform the final  $\text{write}(g)$  opern in schedule  $s'$ .

As can be seen, view equivalence is also based purely on reads & writes alone.

- A schedule  $s$  is view serializable if it is view equivalent to a serial schedule.
- Every conflict serializable schedule is also view serializable.
- Below is a schedule which is view serializable, but not conflict serializable.

$T_3$	$T_4$	$T_6$
$\text{read}(g)$		$\dots$
$\text{write}(g)$	$\text{write}(g)$	$\text{write}(g)$

- Every view serializable schedule i.e. not conflict serializable has blind writes.

vii) The schedule below produces same outcome as the serial schedule  $\langle T_1, T_5 \rangle$  yet is not conflict equivalent or view.

$T_1$	$T_5$
<code>read(A)</code>	
$A = A - 50$	
<code>write(A)</code>	
	<code>read(B)</code>
	$B = B - 10$
	<code>write(B)</code>
<code>read(B)</code>	
$B = B + 50$	
<code>write(B)</code>	
	<code>read(A)</code>
	$A = A + 10$
	<code>write(A)</code>

Determining such equivalence requires analysis

Testing for serializability, Precedence graph

i) consider some schedule of a set of trans.

$T_1, T_2, \dots, T_n$

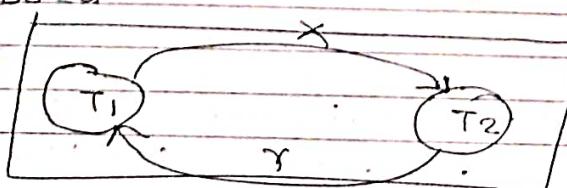
Precedence graph

ii) Precedence graph is a directed graph where the vertices are transaction names

iii) we draw an arc from  $T_i$  to  $T_j$  if two transaction conflict &  $T_i$  accessed the data item on which the conflict arise earlier

iv) we may label the arc by the item that was accessed

e.g

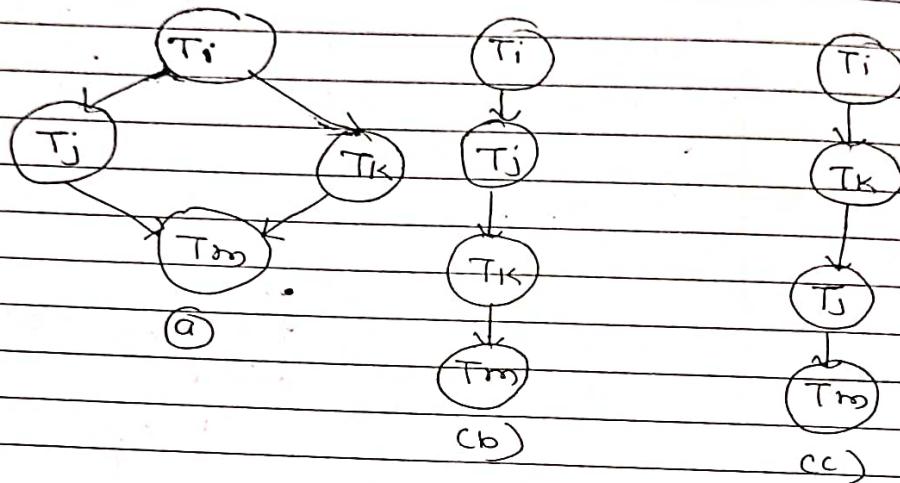


Example ↗

(iv) This is a linear order consistent with the partial order of the graph.

e.g. → serializability order for schedule would be

$$T_5 \rightarrow T_1 \rightarrow T_8 \rightarrow T_2 \rightarrow T_4$$



- A schedule is conflict serializable if & only if its precedence graph is acyclic.
- The precedence graph test for conflict serializability cannot be used directly to test for view -
- Extension to test for view serializability has cost exponential in the size of precedence graph.
- The problem of checking if schedule is view serializable falls in the class of NP complete problems.
- So it checks some sufficient conditions for view serializability, can still be used.

### Recoverable schedules

- If a transaction  $T_i$  reads a data item previously written by trans.  $T_j$ , then the commit operation of  $T_i$  appears before the commit operation of  $T_j$ .
- The following schedule is not recoverable if  $T_2$  commits immediately after the read.

T <sub>8</sub>	T <sub>9</sub>
+read(A)	
write(A)	
	+read(A)
	+read(B)

iii) If T<sub>8</sub> should abort, T<sub>9</sub> would have read an inconsistent DB state, hence db must ensure that schedules are recoverable.

### 2) Cascading Rollbacks

- i) A single transaction failure leads to a series of transaction rollbacks.
- ii) consider the following schedule where none of the transactions has yet committed

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
+read(A)		
+read(B)	.	.
write(A)	+read(A)	
	write(A)	
		+read(A)

- iii) If T<sub>1</sub> fails, T<sub>2</sub>, T<sub>3</sub> must also rolled back.  
It can lead to the undoing of a significant amount of work.

### 2) cascadeless schedules

- i) In cascadeless schedules, cascading rollbacks cannot occur, for each pair of transactions T<sub>i</sub> & T<sub>j</sub> such that T<sub>j</sub> reads a data item previously written by T<sub>i</sub>, the commit operation of T<sub>i</sub> appears before the read operation of T<sub>j</sub>.
- ii) Every cascadeless schedule is also recoverable.
- iii) It is desirable to restrict the schedules to those that are cascadeless

### 3) Strict schedule

- i) If schedule contains no read or write before commit then it is known as strict schedule.
- ii) In this schedule no read/write or write conflict arises before commit hence its strict schedule.

T <sub>1</sub>	T <sub>2</sub>	T <sub>3</sub>
Read(x)		
Read(z)	Read(x)	
Write(x) commit		Read(x) Read(y)
		Write(y) commit
	Read(y)	
	Write(z)	
	Write(y)	
	Commit	

## Lock based protocols

i) Lock is a mechanism to control concurrent access to a data item.

ii) Data items can be locked in two modes

a) Exclusive.

b) Shared.

a) Exclusive (X)

- If transaction  $T_i$  obtained an exclusive-mode lock on item  $Q$ , then  $T_i$  can both read & write  $Q$ .
- It is denoted by  $X$ .

b) Shared

- If transaction  $T_i$  is obtained a shared mode lock on item  $Q$  then  $T_i$  can read but cannot write  $Q$ .
- It is denoted by  $S$ .

- Lock requests are made to the concurrency control manager by the programmer. Transaction can proceed only after request is granted.

\* Lock compatibility matrix

	S	X
S	True	false
X	false	false

i) A transaction may be granted a lock on an item if the requested lock is compatible with locks already held on the item by other transactions.

ii) Any no. of transactions can hold shared locks on an item.

iii) But if any transaction holds an exclusive on the item, no other transaction may hold any lock on the item.

iv) If a lock cannot be granted, the requesting transaction is made to wait till all incompatible

locks held by other transactions have been released  
The lock is then directly granted.

v) example -

```
T2 : lock S(A);  
      read(A);  
      unlock(A);  
      lock S(B);  
      read(B);  
      unlock(B);  
      display(A+B)
```

vi) Locking as above is not sufficient to guarantee serializability, if A & B get updated in between the read of A & B, the displayed sum would be wrong.

vii) A locking protocol is set of rules followed by all transactions while requesting & releasing locks. Locking protocols restrict the set of possible schedules.

### Pitfalls of lock based protocols

i) consider the partial schedule.

T <sub>3</sub>	T <sub>4</sub>
lock X(CB)	
read(CB)	
B := B - 60	
write(CB)	
	lock S(A)
	read(A)
	lock S(B)
lock X(A)	

ii) Neither T<sub>3</sub> nor T<sub>4</sub> can make progress executing lock S(B) causes T<sub>4</sub> to wait for T<sub>3</sub> to release its lock on B while executing lock X(A) causes T<sub>3</sub> to wait for T<sub>4</sub> to release its lock on A

such situation is called as deadlock.

- iii) To handle a deadlock one of  $T_3$  or  $T_4$  must be rolled back & its locks released.
- iv) The potential for deadlock exists in most locking protocols, deadlocks are necessary evil.
- v) Starvation is also possible if concurrency control manager is badly designed.  
 e.g. → A trans. may be waiting for an X-lock on an item, while a sequence of other trans. request & are granted an S-lock on same item  
 - The same trans. is repeatedly rolled back due to deadlocks.
- vi) Concurrency control manager can be designed to prevent starvation.

## Two Phase Locking

- i) This is a protocol which ensures conflict serializable schedules.
- ii) It is a concurrency control method which divides the execution phase of trans. into 2 parts.
- iii) If read & write operations introduce the 2nd unlock operation in the transaction then it is said to be two-phase locking protocol.

It can be divided into two phases

- 1) Growing phase
- 2) Shrinking phase.

### Growing phase

- A transaction obtains locks but may not release any lock.

### Shrinking phase

- Transaction may release locks, but may not obtain any lock.

- Two phase locking does not ensure freedom from deadlocks.
- cascading roll backs is possible under two phase locking so avoid this, follow modified protocol called strict two 2PC

### Types of 2 phase locking protocol

- i) Strict two phase locking protocol.
  - i) It avoids cascading rollbacks.
  - ii) This protocol not only requires two phase locking but also all exclusive locks should be held until transaction commits or aborts.
  - iii) It is not deadlock free.
  - iv) It ensures that if data is being modified by one transaction, then other transaction cannot read it until first transaction commits.
  - v) most DB system implement rigorous two phase locking protocol.
- 2) Rigorous 2 phase locking.
  - i) It avoids cascading rollbacks.
  - ii) It requires that all the shared & exclusive locks to be held until the transaction commits.

## Implementation of Locking

- i) A lock manager can be implemented as a separate process to which transactions send lock & unlock requests.
- ii) The lock manager replies to a lock request by sending a lock grant msg.
- iii) The requesting transaction waits until its request is answered.
- iv) It maintains data structure called a lock table to record granted locks & pending requests.
- v) The lock table is usually implemented as an in memory hash table indexed on the name of data item being locked.

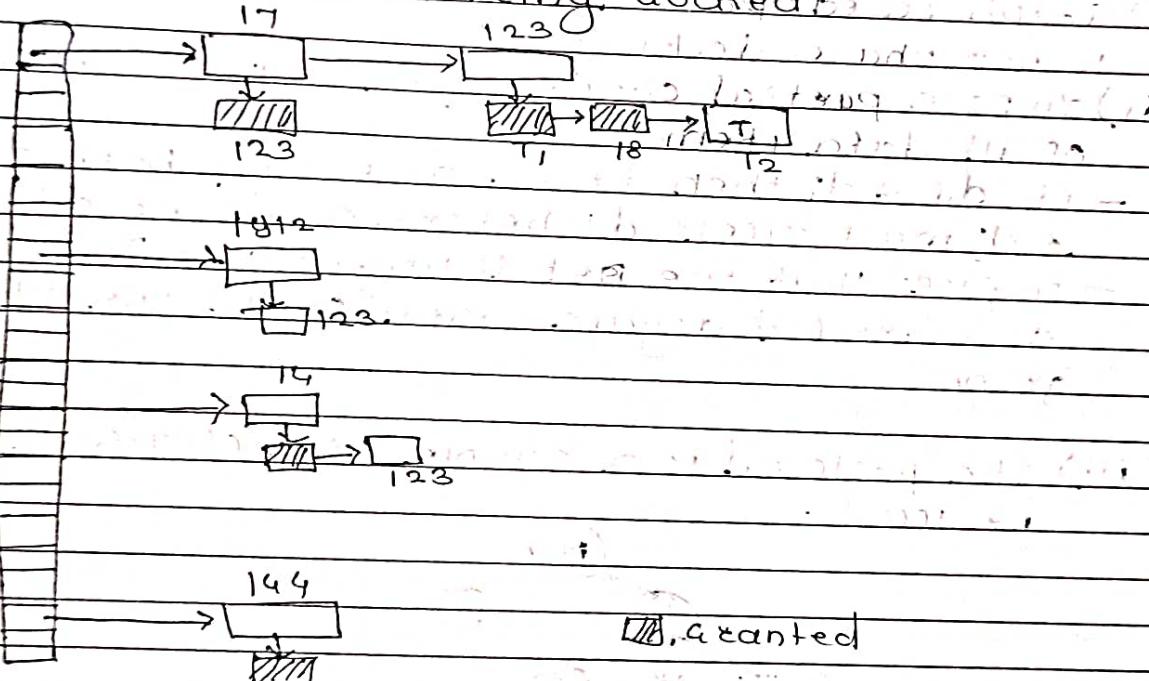


fig. lock table .

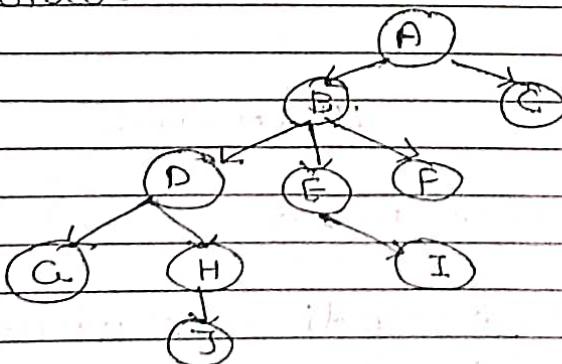
- Dashed rectangles indicates granted locks white one indicates waiting requests.
- Lock table also records the type of lock granted ex. requested.

- New requests result in the request being deleted & later, request been are checked to see if they can now be granted.
- If +Unlock requests result in the request being deleted & later, response requests are checked to see if they can now be granted.
- If transactions about cell waiting or granted requests of the transaction are deleted.

vi) lock manager keep a list of locks held by each transaction to implement this efficiently.

### \* Graph Based Protocol

- i) Graph based protocols are an alternative to two phase locking.
- ii) Impose partial ordering on set  $D = (d_1, d_2, \dots, d_n)$  of all data items.
  - if  $d_i \rightarrow d_j$  then trans. accessing both  $d_i$  &  $d_j$  must access  $d_i$  before accessing  $d_j$ .
  - Implies that the set  $D$  may now be viewed as a directed acyclic graph called the DB graph.
- iii) The protocol is a simple kind of graph protocol.



- a) only exclusive locks are allowed.
  - b) The first lock by  $T_i$  may be on any data item, subsequently a data item can be locked by  $T_i$  only if the present of  $Q$  is currently locked by  $T_i$ .
  - c) Data items may be unlocked at any time.
  - d) A data item that has been locked & unlocked by  $T_i$  cannot subsequently be redlocked by  $T_i$ .
- iv) The three protocol ensures conflict serializability as well as freedom from deadlock.
- v) Unlocking may occur earlier in the three locking protocol than in two phase locking protocol.
  - waiting time is short & concurrency increased
  - protocol is deadlock free, no rollback required

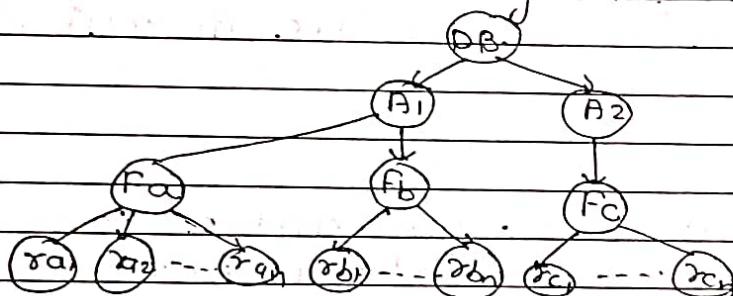
#### vi) Drawbacks

- i) Protocol does not guarantee recoverability or cascade freedom. Need to introduce commit dependencies to ensure recoverability.
- ii) Trans. may have to lock data items that they don't access.

## Multiple Granularity

- i) Allow data items to be of various sizes & define a hierarchy of data granularities, where the small granularities are nested within large ones.
- ii) It can be represented graphically as tree.
- iii) When transaction locks a node in the tree explicitly it implicitly locks all the nodes descendents in the same mode.
- iv) Granularity of locking
  - a) fine granularity → lower in tree, high concurrency, high locking overhead.
  - b) coarse granularity → higher in tree, low locking overhead, low concurrency.

## Example + Granularity Hierarchy



## Deadlock Handling

- consider the following two transactions
- |                  |                  |
|------------------|------------------|
| $T_1$ : write(X) | $T_2$ : write(Y) |
| write(Y)         | write(X)         |

- schedule with deadlock

$T_1$	$T_2$
Lock X on X	
write(X)	

$T_1$	$T_2$
	Lock X on Y
	write(X)

wait for lock X on X

wait for lock X on Y

i) System is deadlocked if there is set of transaction in the set is waiting for another transaction in the set.

ii) Deadlock prevention protocols ensure that the system will never enter into deadlock state.

Some prevention strategies :

a) - Require that each transaction locks all its data items before it begins execution

b)